

# DA5401 Assignment #3

---

**Roll no : DA24C021**

**Name : Venkatesh Duraiarasan**

## **DA5401 Assignment #3**

### **Task 1: Baseline OLS Model and Initial Evaluation**

Objective:

Baseline Model:

Model Summary:

Inference:

### **Task 2: Exploratory Data Analysis (EDA) and Feature Selection**

Task Overview

Visualization and Initial Inferences

Correlation Analysis

Key Insights from Correlation:

Inferences and Recommendations

Conclusion

### **Task 3: Fitting OLS and Reducing Loss**

Objective:

Standardising Input Variables:

Baseline Model:

Model Refinement Process:

Model Evaluation:

Inferences:

### **Task 4 : Regression Model Comparison Using LazyPredict**

Task Overview

Data Preparation

Building and Evaluating Models with LazyPredict

Comparison with OLS Regression

Inference:

# Task 1: Baseline OLS Model and Initial Evaluation

---

## Objective:

The task was to fit an Ordinary Least Squares (OLS) model on the dataset and evaluate the baseline Residual Sum of Squares (RSS) to establish a starting point for further model improvement.

## Baseline Model:

- **Formula:**  $y = m_0 + m_1x_1 + m_2x_2 + m_3x_3 + m_4x_4 + m_5x_5$
- **Residual Sum of Squares (RSS):** 71,877.84

## Model Summary:

- **R-squared:** 0.999, indicating that the model explains nearly all the variance in the data.
- **Coefficients:**
  - Intercept: -9655.31
  - x1: -1067.37
  - x2: 0.10
  - x3: -0.06
  - x4: 284.36
  - x5: 1.63
- **Condition Number:** 1.23e+05, indicating potential multicollinearity among the predictors.

## Inference:

- Despite a high R-squared value, the model's RSS is quite large, suggesting poor fit in absolute terms.
- The large condition number indicates strong multicollinearity, meaning that some of the predictors may be linearly dependent.

# Task 2: Exploratory Data Analysis (EDA) and Feature Selection

---

## Task Overview


The goal of Task 2 was to conduct Exploratory Data Analysis (EDA) on a dataset to understand the relationships between predictor features and how they influence each other, as well as their impact on the output variable, **y**.

## Visualization and Initial Inferences










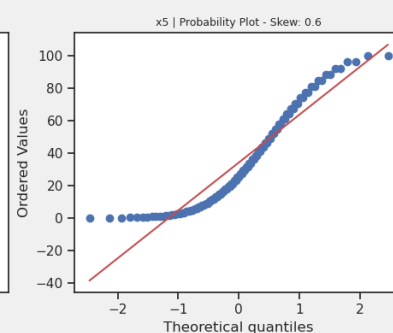


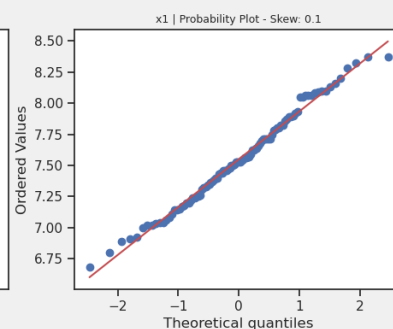


### Violin Plot & Distribution Analysis:


- **x1**: Symmetric distribution with very slight skewness ( $\sim 0.1$ ), closely follows a normal distribution.
- **x2**: Symmetric but with fat tails, indicating more extreme values compared to a normal distribution.
- **x3**: Bimodal and left-tailed, suggesting the presence of two distinct groups within the data.
- **x4**: Symmetric and normal distribution.
- **x5**: Right-skewed with a skewness of  $\sim 0.6$ , and fat tails, indicating higher values are more frequent than expected.
- **y**: Symmetric and normal, with very little skewness ( $\sim 0.3$ ).

### Violin Plot

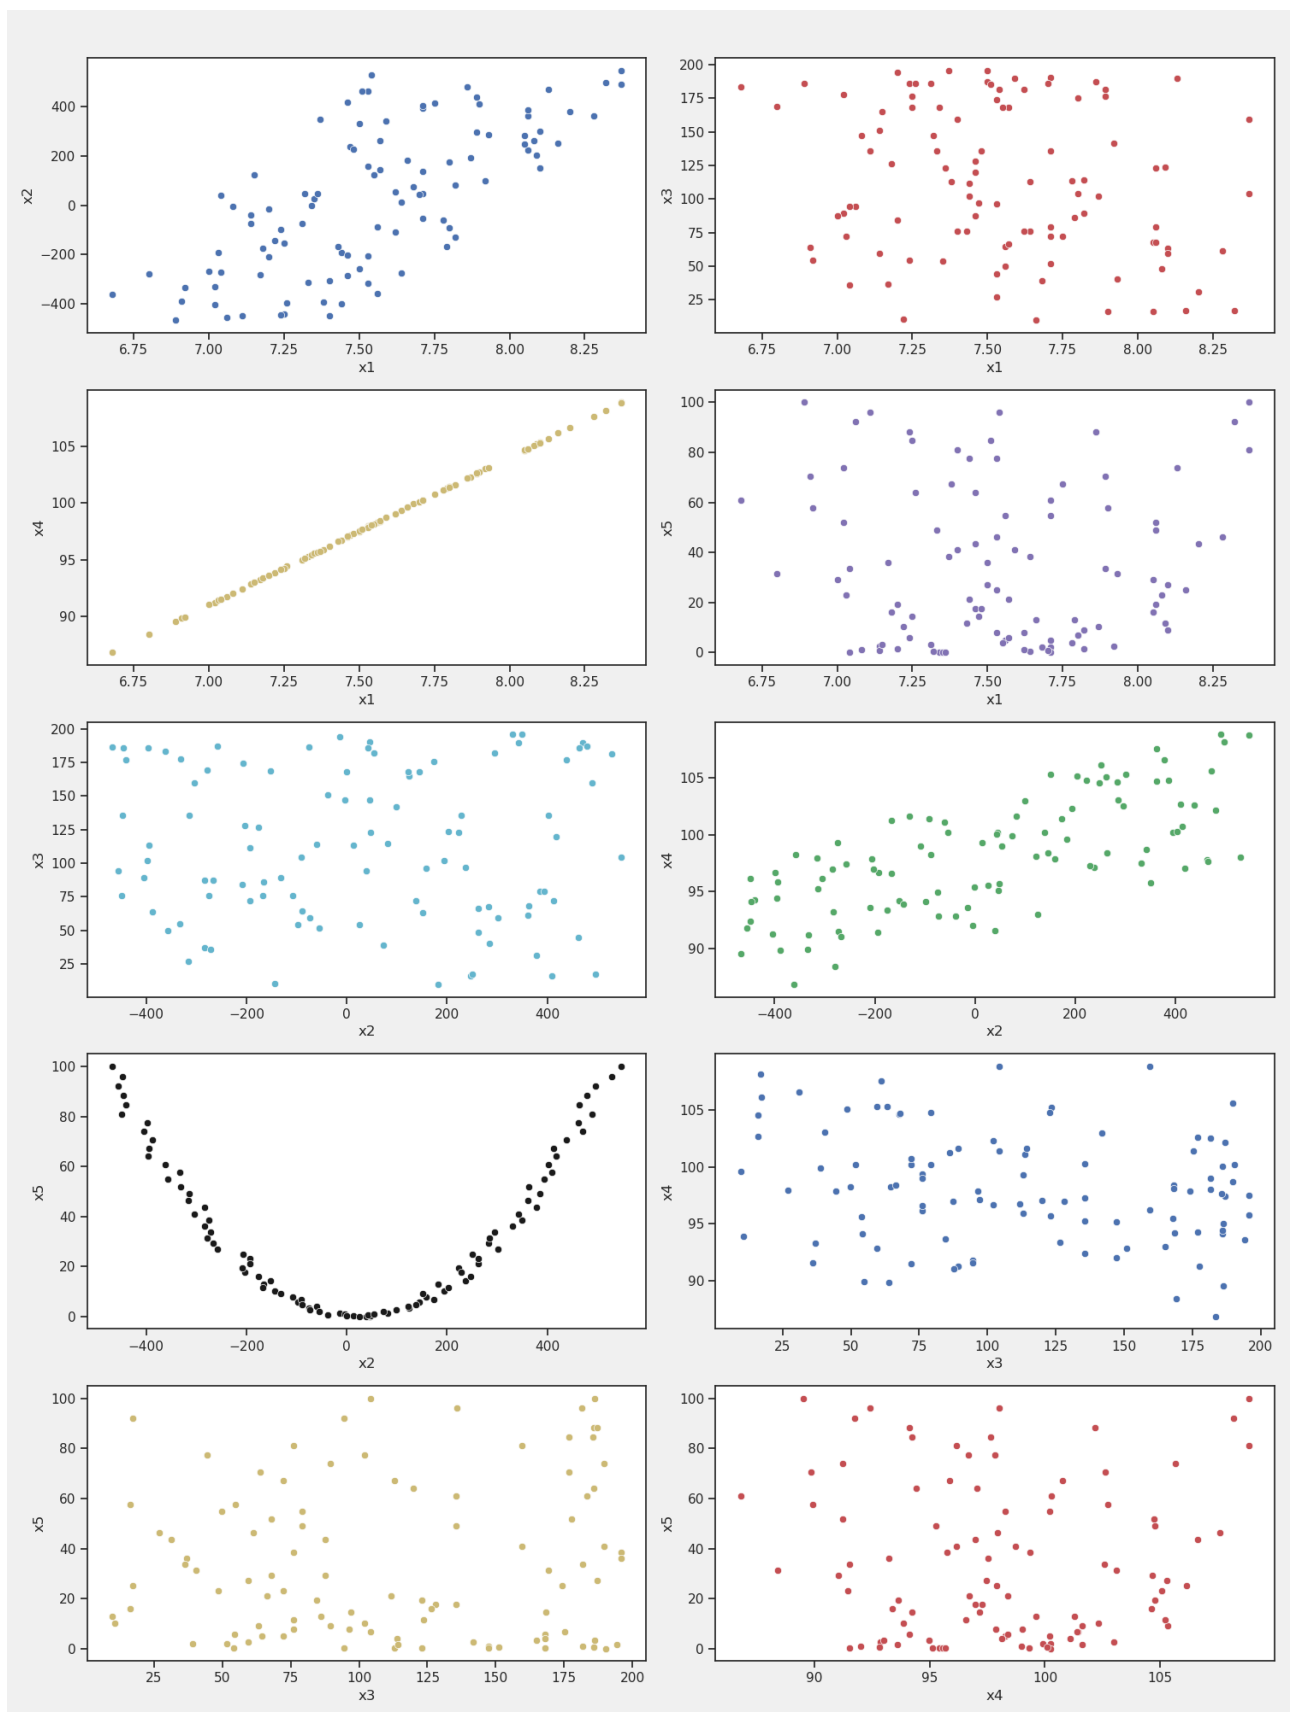


### Distribution Plot, Boxplot, QQ plot






**Pair Plot:**

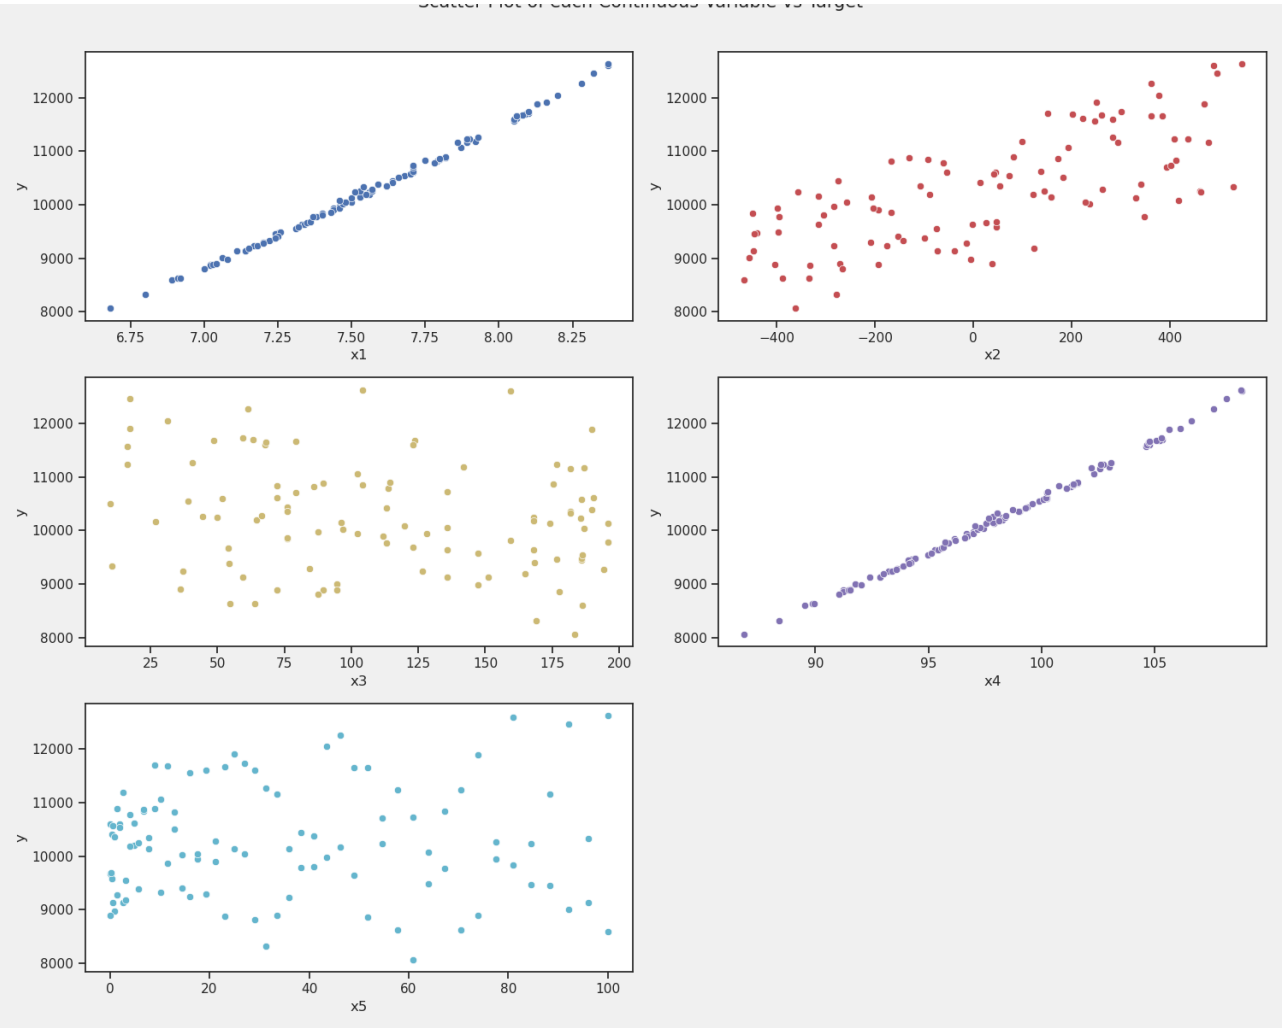


## Correlation Analysis

### Heat Map & Correlation Matrix



Scatter\_Plots



## Key Insights from Correlation:

- **x1 and x4:** Nearly perfect correlation (0.999), indicating that one of these variables is redundant. Since they contribute similar information, one can be removed from the model.
- **x1 and y:** High positive correlation (0.998), suggesting that **x1** is a strong predictor of **y**.
- **x2 and x5:** Slight positive correlation (~0.003), but the scatter plot indicates a potential quadratic relationship. This relationship is weaker compared to others.
- **x3:** Shows a weak correlation with **y**, indicating it might have less predictive power.

## Inferences and Recommendations

- **Remove x4 or x1:** Due to the very high correlation with **x1**, **x4** any one variable can be safely removed without loss of information.
- **Remove x5:** The weak and non-linear relationship with **y**, combined with its quadratic relationship with **x2**, suggests that **x5** may be unnecessary.
- **Consider transformations:** If **x3** is to be included, further exploration of non-linear modeling techniques or transformations (e.g., logarithmic or polynomial) might improve its relationship with **y**.

## Conclusion

Based on the EDA, it is recommended to exclude **x4** and **x5** from the modeling process. The remaining features, particularly **x1** and **x2**, should be considered for modeling.

## Task 3: Fitting OLS and Reducing Loss

---

### Objective:

The goal was to fit an Ordinary Least Squares (OLS) model using selected and transformed features, and to check if the model's loss (measured by the Residual Sum of Squares, RSS) was reduced compared to the baseline model.



## Standardising Input Variables:

Explanatory variables  $X$  are standardised with mean 0 and variance 1, so that corresponding weight terms can be made comparable.

### Baseline Model:

- **Formula:**  $y = m_0 + m_1x_1 + m_2x_2 + m_3x_3 + m_4x_4 + m_5x_5$
- **RSS: 71,877.84** - remains same after scaling input variable

### Model Refinement Process:

#### 1. Removing Features:

- **Attempt 1 & 2:** Removing  $x_2$  or  $x_5$  led to a significant increase in RSS (~112,079), indicating worse performance.
- **Attempt 3 & 4:** Removing  $x_1$  or  $x_4$  showed a slight improvement, but RSS remained higher than the baseline.

#### 2. Adding Interaction Terms:

- **Attempt 5:** Adding the interaction term  $I(x_2 \times x_5)$  slightly improved the RSS to 71,875.46.
- **Attempt 6:** Adding the interaction term  $I(x_1 \times x_4)$  significantly reduced the RSS to **43.30**.
- **Attempt 7:** Combining both interaction terms ( $I(x_1 \times x_4)$  and  $I(x_2 \times x_5)$ ) further reduced the RSS to 42.37.

#### 3. Final Model:

- **Attempt 8:** The final model included  $x_2$ ,  $x_4$ ,  $x_5$ , and the interaction term  $I(x_1 \times x_4)$ . This model achieved an **RSS of 44.63** with an **MSE of 0.0661** on the training data.

$$y = m_0 + m_1x_2 + m_2x_4 + m_3x_5 + m_4x_1x_4$$

### Model Evaluation:

- **Training MSE: 0.4774**
- **Testing MSE: 0.3573**
- **Condition No. : 3.20**

OLS Regression Results						
=====						
Dep. Variable:	y	R-squared:	1.000			
Model:	OLS	Adj. R-squared:	1.000			
Method:	Least Squares	F-statistic:	5.625e+07			
Date:	Sun, 25 Aug 2024	Prob (F-statistic):	8.59e-305			
Time:	10:34:21	Log-Likelihood:	-102.07			
No. Observations:	101	AIC:	214.1			
Df Residuals:	96	BIC:	227.2			
Df Model:	4					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
Intercept	1.022e+04	0.091	1.12e+05	0.000	1.02e+04	1.02e+04
x2	29.0633	0.095	304.566	0.000	28.874	29.253
x4	992.4180	0.096	1.04e+04	0.000	992.228	992.608
x5	39.4991	0.072	548.433	0.000	39.356	39.642
I(x1 * x4)	24.1854	0.061	398.386	0.000	24.065	24.306
=====						

## Inferences:

The introduction of **interaction terms** significantly improved the model's performance, reducing the loss from the baseline RSS of 71,877.84 to **44.63**

Refined model also **removed multi-collinearity issues** in input data as indicated by the reduction in condition number from  $1.23\text{e}+05$  to just **3.20**, in the final model

Model has a **high intercept** which indicates y, the output value, has an offset

Transformed  $x_4$  has higher weight than other input variables

$R^2 = 1$  indicates all the variance in the output is captured by modified features

This refined model demonstrated good generalization with a low MSE(**0.357**) on both training and testing datasets.

# Task 4 : Regression Model Comparison Using LazyPredict

---

## Task Overview

In Task 4, the goal was to use the `lazypredict` package, specifically the `LazyRegressor` class, to automatically build and evaluate various regression models. The performance of these models, as measured by their Root Mean Squared Error (RMSE), was compared against the RMSE of an Ordinary Least Squares (OLS) regression model.

## Data Preparation

The dataset was split into training and testing sets to evaluate the models

## Building and Evaluating Models with LazyPredict

Using `LazyRegressor`, various regression models were automatically built and evaluated.

The output from `LazyRegressor` included a wide range of regression models, each with their respective RMSE,  $R^2$  score

Model	Adjusted R-Squared	R-Squared	RMSE
Lasso	1.00	1.00	43.02
LassoCV	1.00	1.00	43.03
RidgeCV	1.00	1.00	43.26
SGDRegressor	1.00	1.00	43.29
LassoLars	1.00	1.00	43.48
BayesianRidge	1.00	1.00	43.72
Ridge	1.00	1.00	43.73
OrthogonalMatchingPursuitCV	1.00	1.00	44.73
LassoLarsIC	1.00	1.00	44.73
LassoLarsCV	1.00	1.00	44.73
Lars	1.00	1.00	44.73
LinearRegression	1.00	1.00	44.73
RANSACRegressor	1.00	1.00	44.73
TransformedTargetRegressor	1.00	1.00	44.73
HuberRegressor	1.00	1.00	46.02
PoissonRegressor	1.00	1.00	46.24
PassiveAggressiveRegressor	1.00	1.00	46.99
OrthogonalMatchingPursuit	1.00	1.00	70.15
LarsCV	0.99	0.99	94.60
ExtraTreesRegressor	0.99	0.99	138.95
GradientBoostingRegressor	0.98	0.99	148.53
BaggingRegressor	0.98	0.99	157.87
XGBRegressor	0.98	0.99	158.89
RandomForestRegressor	0.98	0.98	162.51
AdaBoostRegressor	0.97	0.98	185.88
DecisionTreeRegressor	0.97	0.98	194.79

## Comparison with OLS Regression

For comparison, an Ordinary Least Squares (OLS) regression model was from Task 3 and its RMSE was calculated to be **0.357**

## Inference:

Refined Linear regression model performed better than above general models due to proper feature selection and feature transformation

Also different techniques report different performance metrics due to

- overfitting,
- model complexity (ability to capture non-linear relationships in the data more effectively)